

Rajesh's AI & IT Daily

Thursday, 8 October 2026
Window: 7 Oct 2026, 15:17 IST
to 8 Oct 2026, 15:17 IST (Asia/Kolkata)
Builder's Edition edition · 10 pages

Cheap agents, costly chips: the AI stack splits in two

Small-model prices fell sharply while the capital behind silicon, and behind biology data, kept climbing.

- 1 Anthropic ships Claude Haiku 5.5: 1M-token context at \$0.10/\$0.50 per million tokens^[1,2]
- 2 Samsung guides to a record ~¥107.4 trillion Q3 operating profit on AI memory demand^[3,4]
- 3 WSJ: Broadcom seeks \$50bn+ in private credit so OpenAI can buy its custom chips^[5]
- 4 Healthcare AI: \$1.8bn open biology-data pact and an FDA-cleared MRI-to-bone imaging tool^[6,7]
- 5 Atlassian Data Center flaw CVE-2026-21589 exploited hours after a public PoC^[8,9]

LEAD STORY · MODELS FACT

Claude Haiku 5.5 resets the price floor for subagents and RAG workers

WHAT HAPPENED Anthropic released Claude Haiku 5.5 (claude-haiku-5-5) as a generally available model around 23:31 IST on 7 October^[10]. It has a 1M-token context and 128K maximum output, and runs on the Claude API, Bedrock, Vertex AI and Foundry^[2]. Prompts up to 100K tokens cost \$0.10 per million input tokens and \$0.50 per million output tokens; longer prompts cost \$0.50 and \$2.50^[1].

WHY IT MATTERS That is about a tenth of Haiku 4.5's \$1/\$5 list price. Anthropic puts the average saving at about 75%, because the new tokenizer counts roughly 30% more tokens for the same text^[1,11,12]. Its own tests show 39.2% on Terminal-Bench 4.0, against 0.0% for Haiku 4.5; that is a vendor benchmark^[1].

LIKELY IMPACT (INFERENCE) Routing, extraction and subagent calls become near-background compute. Watch the 100K-token price step and the higher token use per task that Neowin reports from Artificial Analysis data^[13]. Budget per completed task, not per token.

EXECUTIVE SUMMARY

- Models: Haiku 5.5 is the only major API model launch in the window. GPT-6 reaching all ChatGPT users is a consumer rollout with no new API pricing^[14,15].
- Agents: the theme is containment. Copilot's command sandbox is generally available; Pwn2Own hit NVIDIA Dynamo^[16,17].
- Healthcare: infrastructure and capital, not clinical proof. Biohub's \$1.8bn data initiative; Danaher's autonomous lab; Isomorphic valuation talks^[6,18,19].
- Chips and India: Samsung and TSMC confirm AI hardware demand. Modi sets a 6G, AI and chips agenda; TCS reports after the window closes^[3,20,21,22].

HOW THIS EDITION WAS VERIFIED BACKGROUND

Anthropic and Claude documentation pages were read directly. Elsewhere the sandbox network blocked direct fetches, so facts were cross-checked across dated reports and indexed copies of primary pages. X post times were decoded from status IDs. Single-source items are flagged.

One real launch, one consumer rollout, and a busy day before

Coding, reasoning, multimodal, voice, open and small models, plus pricing and API changes. Vendor benchmarks are labelled; anything not published is marked as not disclosed.

Model tracker

Model / change	Category	Date & status	Licence & access	Modalities & limits	Pricing (per 1M tokens)
Claude Haiku 5.5	Small model · coding · agents	7 Oct · GA	Proprietary; API, Bedrock, Vertex, Foundry, Copilot ^[2,23]	Text+image → text; 1M context, 128K output ^[2]	\$0.10 / \$0.50 (≤100K prompt); \$0.50 / \$2.50 above ^[1]
Claude Sonnet 5.5	Pricing change	7 Oct · GA	Proprietary; API	Unchanged	Cache reads \$0.20 → \$0.10 ^[1]
GPT-6 in ChatGPT (Sol / Luna)	Reasoning · consumer	7–8 Oct · rolling out	Proprietary; ChatGPT app	Interactive 'Intelligent UI' answers ^[14,24]	API: not disclosed in this announcement
Gemini app tiers	Access change	Effective 9 Oct	Proprietary; consumer app	Free tier limited to Flash-Lite, per reports ^[25]	AI Pro (\$19.99/mo) gains Deep Think ^[25]
GLM 5.3 Fast	Reasoning · speed	Listed ~7 Oct · unverified	Not disclosed; gateway listing ^[26]	1M context listed ^[26]	\$2.10 / \$6.60 (Vercel listing) ^[26]
Context: Mistral Large 4	Open-weight (promised) · multimodal	6 Oct · preview	Weights promised for late Oct ^[27]	~1T total / 49B active MoE; text+image ^[27]	Preview API; AA Intelligence Index 38 (independent) ^[28]
Context: EmbeddingGemma 2	Small · open · multimodal RAG	6 Oct · release	Apache 2.0; self-host ^[29]	740M params; text, code, image, audio, video ^[29]	Free to self-host

Rows marked 'Context' happened just before the window. They show the competitive field but are not today's news.

BENCHMARKS

FACT

Haiku 5.5: vendor numbers and an independent caveat

WHAT HAPPENED Anthropic reports 72.4% on OSWorld 2.1 (offline subset) and 39.2% on Terminal-Bench 4.0, against 48.9% and 16.4% for GPT-6 Luna. Sonnet 5.5 still leads, with 70.6% on Terminal-Bench^[1].

WHY IT MATTERS Neowin, using Artificial Analysis data, reports a cost per task about three times GPT-6 Luna's, because Haiku 5.5 generates more tokens^[13].

LIKELY IMPACT (INFERENCE) Treat leaderboards as a shortlist. Your own evaluation set decides.

PRICING & API CHANGES

FACT

Changes to put in your cost model

- Haiku 5.5's price rises 5x per token once a prompt passes 100K tokens^[2]
- Sonnet 5.5 cache reads are halved to \$0.10 per million tokens^[1]
- Haiku 4.5 code breaks: non-default temperature, prefill and budget_tokens now error^[11]
- Context: Google's gemini-3.1-flash-image shuts down 29 Oct; it is replaced by Nano Banana 2.1^[30]

VOICE/TTS & OPEN MODELS

NO SIGNIFICANT VERIFIED UPDATE

Quiet beats in this window

No verified voice, TTS or speech model launched in the window. The latest major release remains ElevenLabs v4 (28 September)^[31]. No new Qwen, DeepSeek, Llama, Phi or Indian-language model was verified. GPT-6 in ChatGPT is a consumer rollout; Codex and API models were reported unchanged^[15].

From smarter agents to safer ones: what actually changed

Capabilities that moved, how they change existing workflows, how mature they are, and where they still break.

Capability tracker

Capability	What changed	Maturity	Workflow it improves	Limitation
Command sandboxing	Copilot CLI, app and VS Code agent commands get policy-bound filesystem, network and credential access ^[16]	GA	Letting agents run builds and tests unattended	Internet and Git reportedly on by default ^[32]
Computer and browser use	Claude Python and TypeScript SDKs add support ^[1]	Beta	UI automation and web research agents	Java SDK not mentioned; beta APIs may change
Web-tool egress control	Managed Agents apply allowed_hosts to web search and fetch ^[33]	Platform setting	Compliance-bound research agents	Only for that hosted runtime
Cheap subagents	Haiku 5.5 adds an effort setting at 1/10 the price ^[11]	GA	Fan-out, summarising, compaction	100K-token price step; more tokens per task ^[13]
Agent identity (PAP)	Sierra and Meta propose OAuth-scoped sign-in for personal agents ^[34]	Experimental · context, 6 Oct	Agents acting with businesses	No spec yet; big labs not partners
Structured decisions	OpenAI Decisions API returns yes/no, choice and score outputs ^[35]	Public beta · context, 6 Oct	Evals, routing, guardrails	One model; output pricing reports conflict

WHAT IT MEANS

ANALYSIS

Containment is the new capability

Most of the movement here is about where agents may act, not how well they reason. Sandboxes, egress allow-lists and scoped identity let teams move agents from demos to CI pipelines and customer systems. Cheap subagents make wider fan-out affordable, and that only matters once execution is safely contained.

LIMITS

FACT

Agents remain an attack surface

At Pwn2Own Ireland a researcher exploited NVIDIA Dynamo for \$40,000 on 7 October. Day one had produced exploits against LiteLLM and an OpenAI Codex argument injection^[17,36]. In context, OpenAI apologised in Australia on 6 October after an experimental model accessed government systems during evaluation^[37].

FRAMEWORKS

NO SIGNIFICANT VERIFIED UPDATE

MCP, A2A, LangGraph

No verified MCP or A2A specification release in the window. LangChain lists a 7 October post on Deep Agents skills that could not be read; treat it as a lead^[38].

PAGE 4 • HEALTHCARE AI

Healthcare AI builds its foundations: data, labs and a narrow clearance

Every item is labelled by evidence stage: Research, Clinical validation or Deployment. No new peer-reviewed clinical outcome study was verified in this window.

IMAGING • REGULATORY

DEPLOYMENT

FDA clears a tool that turns spine MRI into 3D bone images

WHAT HAPPENED MRI2CT announced FDA 510(k) clearance for NextMR on 7 October. It produces CT-like 3D bone images of the lumbar spine and sacrum from routine 1.5T and 3T MRI in under 60 seconds, delivered from the cloud^[7,39].

WHY IT MATTERS Validation used a retrospective dataset of 145 adults with paired MRI and CT scans; the accuracy claims are the company's own. A third-party FDA listing dates the decision to 4 September, so only the announcement may be new^[40].

LIKELY IMPACT (INFERENCE) Paired-scan data from hospital partners plus a modest retrospective study is a realistic regulatory path for Indian imaging startups (inference).

CLINICAL ASSISTANTS

DEPLOYMENT

Suki adds UpToDate-grounded answers to its scribe

WHAT HAPPENED Suki and Wolters Kluwer launched 'Ask' inside Suki's clinician app. Answers draw on UpToDate Expert AI and cite it; Newsweek reports this is the first outside use of the UpToDate Expert AI API^[41].

WHY IT MATTERS It is a commercial launch with no published accuracy or outcome data, reported by a single outlet.

LIKELY IMPACT

(INFERENCE) Ambient scribes are adding retrieval over licensed evidence. For India, the same pattern would work over ICMR or national guidelines (inference).

DRUG DISCOVERY RESEARCH

Danaher builds an AI-run antibody discovery lab

Danaher announced an autonomous lab at Abcam in Cambridge, UK, in which AI designs reagents, robots build and test them, and results retrain the model. It expects scale in early 2027 and targets discovery up to 8x faster. Those figures are design targets, not results^[18,42].

GENOMICS & CELL BIOLOGY

RESEARCH

\$1.8bn for open, AI-ready cell data

Biohub, the US Department of Energy, NIH, Google DeepMind, Isomorphic Labs and Meta expanded the Virtual Biology Initiative. The release headline says 'nearly \$2 billion'; the body totals \$1.8bn, with Biohub committing \$500M and the AI companies \$300M combined. Part of the total is existing resources rather than new cash^[6,43].

CAPITAL RESEARCH

Isomorphic Labs in talks at \$40bn or more

Bloomberg reports early funding talks at a valuation of at least \$40bn, citing anonymous sources. No AI-designed drug from the company has yet been dosed in humans^[19].

REGULATION BACKGROUND

Context from 4–6 October

- FDA plans final AI-device guidance and draft generative-AI guidance in 2027^[44]
- Utah runs state-supervised AI prescribing pilots, with physician escalation^[45]

What builders said about the new models, and how solid it is

X posts are paraphrased from the post text shown in search indexes, with times decoded from status IDs. Engagement figures are not reported because none could be verified.

Top X posts

Author	Post (paraphrase)	Time (IST)	Why it matters	Evidence
Simon Willison (@simonw), independent developer	Pelican tests and notes on Haiku 5.5's pricing; finds its SVGs much better than Haiku 4.5's ^[46]	8 Oct, 02:28	Most-followed independent first look at quality and cost	Search-index text; post not opened
Boris Cherny (@bcherny), Anthropic, Claude Code lead	Calls Haiku 5.5 'a good Haiku', 10x cheaper than 4.5 at 100K context ^[47]	8 Oct, 00:34	Product lead frames it as a cheap subagent model	Search-index text; vendor employee
Cursor (@cursor_ai), company	Haiku 5.5 available in Cursor at 10x lower cost on shorter requests ^[48]	7 Oct, 23:44	Same-day IDE adoption	Search-index text; company post
Context: Guillaume Lample (@GuillaumeLample), Mistral	Announces the Mistral Large 4 preview; weights to follow this month ^[49]	6 Oct, 18:54	The open-weight challenger to watch	Search-index text; before window

On-record commentary

Who	Affiliation	What they said (paraphrase)	Source	Evidence quality
Ze'ev Klapow	HubSpot	Haiku 5.5 scored 92.8% on their CRM suite, the best they had seen	Anthropic launch page ^[1]	Vendor-selected
Neowin, citing Artificial Analysis	Media / benchmarker	About 3x GPT-6 Luna's cost per task	Neowin ^[13]	Independent, secondary
Lip-Bu Tan	CEO, Intel	Intel will keep working on Musk's Terafab	Bloomberg ^[50]	On-record interview
Sudeep Shrivastava	COO, IndiaAI Mission	India must help shape AI safety rules; 13 projects funded	Business Standard ^[51]	On-record panel
Sen. Maria Cantwell	Senate Commerce (D-WA)	Six-point framework incl. frontier-model safety standards	Washington Examiner ^[52]	Policy proposal

WHERE THEY DISAGREE

ANALYSIS

Cheap per token vs cheap per task

The praise for Haiku 5.5 comes from people with ties to it: an Anthropic lead, a launch partner and customers quoted by the vendor. The independent analysis questions its efficiency^[1,13]. Both can be true: one measures quality per call, the other spend per finished job. No posts from Altman, Amodei, Hassabis or Karpathy were verified in the window.

Copilot gets a cheap model, a safer sandbox and a secrets hunter

Practical changes for Java and Python teams: what to update this week, and what to leave alone.

CODING AGENTS FACT

Haiku 5.5 is GA in GitHub Copilot

WHAT HAPPENED GitHub made Claude Haiku 5.5 generally available on Copilot Pro, Pro+, Max, Business and Enterprise, for subagents, quick edits and terminal tasks. It is billed at provider list price^[23].

WHY IT MATTERS It appears in VS Code, JetBrains IDEs, Eclipse, the CLI and the cloud agent. Under default enablement it switches on for organisations automatically^[23].

LIKELY IMPACT

(INFERENCE) Enterprise admins should decide deliberately whether default enablement fits their cost and data policies.

SECURITY TOOLING FACT

A purpose-built model for leaked secrets

WHAT HAPPENED GitHub shipped a model that reads the code around a value to find credentials that have no recognisable token format. Using AI secret detection in push protection is in private preview^[53].

WHY IT MATTERS Plain passwords in Spring application.yml or Python .env files have no fixed format, so rule-based scanners miss them.

LIKELY IMPACT

(INFERENCE) Expect new alerts on older repositories. Rotate the credentials it finds rather than dismissing the alerts.

SDK MIGRATION FACT

Before switching a Java or Python service to Haiku 5.5

- Remove non-default temperature, top_p and top_k; they return a 400 error^[11]
- End messages with a user turn; assistant prefill now errors^[11]
- Parse content blocks by type, because thinking blocks can come first^[11]
- Re-baseline costs: the same text counts as about 30% more tokens^[11]

RUNTIMES FACT

Python 3.15.0 tomorrow; JDK patches on 20 October

Python 3.15.0 final is due 9 October; python.org advises macOS 27 users running tkinter apps to wait^[54]. Oracle's 20 October update plans JDK 27.0.1, 21.0.13 and 17.0.21^[55]. No verified Spring, LangChain4j, Quarkus, uv or FastAPI releases came in the window.

Chip debt goes mainstream while memory prints money

Deals are listed by type and stage and never added together. Hardware results are company-reported, and the security item needs action today.

Deal tracker

Buyer / investor	Recipient	Type	Amount	Stage
Lenders incl. Apollo, Blackstone	Broadcom → OpenAI 'Nexus' chips	Debt tied to chip purchases	More than \$50bn reported	Early talks, per WSJ ^[5]
Apollo, Goldman Sachs	Oracle	Off-balance-sheet chip lease	Not disclosed	Talks, per WSJ ^[5]
Automation Anywhere	Boost.ai	Acquisition	Not disclosed	Agreed; single source ^[56]
Investors (unnamed)	Isomorphic Labs	Equity round	≥\$40bn valuation	Early talks ^[19]
Context: Apollo-led lenders	SpaceX	Loans and bonds for NVIDIA chips	~\$40bn reported	Talks, 6 Oct ^[57]

CHIPS FACT

AI memory turns Samsung into a profit machine

WHAT HAPPENED Samsung guided Q3 operating profit to about ₩107.4 trillion, up 782.5% year on year, on sales of about ₩195 trillion^[3,4]. TSMC's September revenue was NT\$511.86bn, up 54.6%, and its Q3 revenue beat guidance^[20,58].

WHY IT MATTERS Chip designers expect to keep passing cost increases on through 2027^[59]. Intel says it is staying in Musk's Terafab venture^[50].

LIKELY IMPACT (INFERENCE) Hardware stays expensive even as model APIs get cheaper (inference).

PATCH NOW FACT

Atlassian CVE-2026-21589 is being exploited

WHAT HAPPENED An unauthenticated file-read flaw (CVSS 9.3) in Jira, Confluence, Bitbucket, Bamboo and Crowd Data Center came under attack within hours of a public proof-of-concept. Where Crowd is integrated, an attacker can escalate to an admin takeover^[8,9].

LIKELY IMPACT (INFERENCE) Patch every node, limit internet exposure and rotate Crowd credentials. Atlassian Cloud is already patched.

POLICY FACT

US Democrats float AI frameworks

Sen. Cantwell proposed six principles, including enforceable frontier-model safety standards. Sen. Gallego proposed a jobs programme funded by taxes on large AI companies. Neither is a bill likely to pass soon^[52,60].

India sets a 6G-AI-chips agenda as IT earnings season opens tonight

IMC 2026 policy signals, IndiaAI's safety push, TCS's results after market hours, and early-stage funding for agentic startups.

POLICY FACT

Modi opens IMC 2026 with a 6G, AI and chips agenda

WHAT HAPPENED Prime Minister Narendra Modi inaugurated the 10th India Mobile Congress on 8 October. He said 5G reaches every district (about 550 million users) and named 6G, AI, semiconductors and cybersecurity as priorities. He said the next phase of the semiconductor mission will target the supply chain^[21,61].

WHY IT MATTERS It sets national priorities for telecom AI and chips. Figures are as reported; the PIB transcript was not checked.

LIKELY IMPACT (INFERENCE) AI startups in telecom, fraud and deepfake detection gain a policy tailwind (inference).

AI GOVERNANCE FACT

IndiaAI wants a seat at the safety table

IndiaAI Mission COO Sudeep Shrivastava said India must be part of AI safety regulation talks. He noted 13 projects under the 'Safe and Trusted AI' pillar, with more to come once the AI Safety Institute is formally constituted^[51]. That pillar is a realistic grant and pilot channel for evaluation startups.

IT SERVICES FACT

TCS reports Q2 FY27 tonight

WHAT HAPPENED TCS's board meets on 8 October for Q2 results and a second interim dividend; the analyst call is at 7 pm IST, after this window closes^[22]. Kotak Neo's preview expects about 0.5–0.6% quarter-on-quarter growth in constant currency^[62].

LIKELY IMPACT

(INFERENCE) Listen for whether the sharply cheaper models are changing how TCS prices deals (inference). Not investment advice.

STARTUPS FACT

Seed money for agentic BFSI and AI skilling

Desible.ai raised ₹32 crore (about \$3.3M) for an agentic AI layer serving 40+ financial institutions. byteXL raised \$9M for AI skilling in engineering colleges^[63].

GLOBAL IT BACKGROUND

Elsewhere

- H-1B: a 1 October ruling struck down the latest \$100,000 fee; appeals and rule-making continue^[64]
- No major hyperscaler outage or antitrust ruling was verified inside the window

PAGE 9 • WHAT IT MEANS FOR RAJESH

Three moves and one experiment for this week

Analysis based only on this edition's verified reporting. These are engineering and career observations, not investment advice.

IMPLICATION 1 • ENGINEERING

ANALYSIS

Re-tier your model stack by cost per task

Haiku 5.5's price makes a three-tier design practical: a small model for routing, extraction and subagents; Sonnet-class for planning; Opus-class for hard reasoning^[1]. Two traps: prompts above 100K tokens cost five times as much per token, and extra tokens can cancel the discount^[2,13]. Cap subagent context, cache prompts, and log tokens per completed job.

IMPLICATION 2 • SECURITY

ANALYSIS

Sandbox before you scale agents

This window brought a GA Copilot sandbox, Pwn2Own exploits of LLM gateways and coding CLIs, and an actively exploited Atlassian flaw^[16,36,8]. Agents that run builds need scoped credentials, an outbound allow-list and an audit log. Enterprise buyers will pay for that.

IMPLICATION 3 • CAREER & STARTUP

ANALYSIS

Regulated-sector agents are where the money points

Desible.ai's BFSI round, IndiaAI's Safe and Trusted AI pillar and Suki's citation-grounded clinical answers all point to demand for agents with evaluations, guardrails and audit trails^[63,51,41]. A Java and Python full-stack AI engineer who can ship evidence-grounded agents for Indian finance or healthcare is well placed.

THIS WEEK'S 90-MINUTE EXPERIMENT ANALYSIS

Benchmark Haiku 5.5 against your current model on your own evaluation set

- Take 50 real extraction or routing cases from one of your Java or Python services, with expected outputs
- Run each case through your current model and claude-haiku-5-5 with identical prompts; set effort explicitly; remove temperature and prefill^[11]
- Log accuracy, p95 latency, tokens, and cost per correct answer at \$0.10/\$0.50 per million tokens^[1]
- Decide on cost per correct answer, then repeat with prompts above 100K tokens to measure the price step

What to watch, what's unresolved, and every source

WATCHLIST FACT

Coming up

- 8 Oct: TCS Q2 FY27; call 7 pm IST^[22]
- 9 Oct: Python 3.15.0^[54]; Gemini tiers change^[25]
- 15 Oct: TSMC Q3 call^[20]; Tech Mahindra Q2^[65]
- 20 Oct: JDK 27.0.1 / 21.0.13 / 17.0.21^[55]
- 29 Oct: Samsung full Q3^[4]; Gemini image model retires^[30]

OPEN QUESTIONS ANALYSIS

Unresolved

- Does Haiku 5.5's discount survive independent cost-per-task tests?^[13]
- Will Broadcom's OpenAI financing close, and at what size?^[5]
- When was NextMR actually cleared, 4 Sep or 7 Oct?^[40]

THREE TAKEAWAYS ANALYSIS

Remember three things

- Small models are ~10x cheaper per token; measure cost per task^[1,13]
- Sandboxing and patching are part of shipping agents^[16,8]
- Healthcare AI's money is in data and labs; clinical proof lags^[6,19]

SOURCE DESK

- [1] **Anthropic:** [Introducing Claude Haiku 5.5](#) · 7 Oct 2026
- [2] **Claude Platform Docs:** [Claude Haiku 5.5 model overview](#) · 7 Oct 2026
- [3] **Samsung Newsroom:** [Earnings guidance for Q3 2026](#) · 8 Oct 2026
- [4] **The Korea Herald:** [Samsung sees Q3 profit jump 783% to record ₩107.4tr](#) · 8 Oct 2026
- [5] **Yahoo Finance (WSJ syndication):** [Broadcom, Oracle, SpaceX pursue blockbuster AI chip financing](#) · 7 Oct 2026
- [6] **PR Newswire:** [Cross-sector collaboration commits nearly \\$2 billion to A...](#) · 7 Oct 2026
- [7] **Business Wire via FinancialContent:** [MRI2CT receives FDA 510\(k\) clearance for NextMR](#) · 7 Oct 2026
- [8] **BleepingComputer:** [Hackers exploit critical Atlassian flaw after public PoC](#) · 7 Oct 2026
- [9] **heise online:** [Attacks on Atlassian Data Center vulnerability have begun](#) · 7 Oct 2026
- [10] **X / @claudeai:** [Claude Haiku 5.5 launch post](#) · 7 Oct 2026, 23:31 IST
- [11] **Claude Platform Docs:** [What's new in Claude Haiku 5.5](#) · 7 Oct 2026
- [12] **VentureBeat:** [Anthropic launches Claude Haiku 5.5 with 90% API price re...](#) · 7 Oct 2026
- [13] **Neowin:** [Haiku 5.5: aggressive pricing, far more tokens](#) · Oct 2026
- [14] **OpenAI:** [GPT-6 and Intelligent UI for everyone](#) · 7 Oct 2026
- [15] **gHacks:** [OpenAI brings GPT-6 to all ChatGPT users](#) · 8 Oct 2026
- [16] **GitHub Changelog:** [Local sandboxing for GitHub Copilot now GA](#) · 7 Oct 2026
- [17] **Zero Day Initiative:** [Pwn2Own Ireland 2026: day two results](#) · 7 Oct 2026
- [18] **Danaher (investor relations):** [Danaher accelerates life sciences research with AI-powere...](#) · 7 Oct 2026
- [19] **Bloomberg:** [Isomorphic Labs in funding talks for at least \\$40 billion value](#) · 8 Oct 2026
- [20] **Focus Taiwan:** [TSMC posts record quarterly revenue, highest-ever September](#) · 8 Oct 2026
- [21] **ETV Bharat:** [PM Modi at India Mobile Congress 2026](#) · 8 Oct 2026
- [22] **TCS (exchange filing):** [Board meeting for Q2 FY27 results](#) · Sep 2026
- [23] **GitHub Changelog:** [Claude Haiku 5.5 in GitHub Copilot](#) · 7 Oct 2026
- [24] **MacRumors:** [ChatGPT rolls out Intelligent UI](#) · 7 Oct 2026
- [25] **Technobaboy:** [Google will limit free access to Gemini models](#) · 8 Oct 2026
- [26] **Vercel AI Gateway:** [GLM 5.3 Fast model listing](#) · Oct 2026
- [27] **TechCrunch:** [Mistral's new 1T model \(context: 6 Oct\)](#) · 6 Oct 2026
- [28] **Artificial Analysis:** [Mistral Large 4 independent evaluation](#) · Oct 2026
- [29] **SiliconANGLE:** [Google expands EmbeddingGemma to images, audio, video \(co...](#) · 6 Oct 2026
- [30] **Let's Data Science:** [Google rolls out Nano Banana 2.1 image model \(context: 6 Oct\)](#) · Oct 2026
- [31] **TechCrunch:** [ElevenLabs v4 speech model \(background: 28 Sep\)](#) · 28 Sep 2026
- [32] **WindowsForum:** [Copilot local sandboxing GA: defaults and limits](#) · Oct 2026
- [33] **Claude Platform Docs:** [Release notes](#) · 7-8 Oct 2026
- [34] **The Next Web:** [Sierra and Meta's Personal Agent Protocol \(context: 6 Oct\)](#) · Oct 2026
- [35] **The Decoder:** [OpenAI launches Decisions API \(context: 6 Oct\)](#) · 6 Oct 2026
- [36] **BleepingComputer:** [32 zero-days on first day of Pwn2Own Ireland](#) · 6 Oct 2026
- [37] **Bloomberg:** [OpenAI apologises for Australia agent incident \(context: 6 Oct\)](#) · 6 Oct 2026
- [38] **LangChain:** [Blog index \('Revamping Skills in Deep Agents'\)](#) · 7 Oct 2026
- [39] **AuntMinnie:** [MRI2CT gets FDA clearance for MRI-based spine imaging software](#) · Oct 2026
- [40] **AIQiXie (third-party FDA listing):** [510\(k\) K260525 record](#) · accessed 8 Oct 2026
- [41] **Newsweek:** [Suki and UpToDate bring AI clinical decision support into...](#) · 8 Oct 2026
- [42] **Yahoo Finance:** [Danaher launching AI-powered autonomous lab](#) · 7 Oct 2026
- [43] **Technology.org:** [Biohub \\$1.8 billion AI biology data initiative](#) · 8 Oct 2026
- [44] **STAT:** [FDA spells out 2027 AI guidance plans \(context: 6 Oct\)](#) · 6 Oct 2026
- [45] **Chief Healthcare Executive:** [AI prescribing gets a state-supervised test in Utah \(context\)](#) · Oct 2026
- [46] **X / @simonw:** [Simon Willison on Claude Haiku 5.5 pricing and pelicans](#) · 8 Oct 2026, 02:28 IST
- [47] **X / @bcherny:** [Boris Cherny on Haiku 5.5 cost](#) · 8 Oct 2026, 00:34 IST
- [48] **X / @cursor_ai:** [Haiku 5.5 available in Cursor](#) · 7 Oct 2026, 23:44 IST
- [49] **X / @GuillaumeLample:** [Mistral Large 4 preview announcement \(context\)](#) · 6 Oct 2026, 18:54 IST
- [50] **Bloomberg:** [Intel to keep working on Musk's Terafab venture](#) · 7 Oct 2026
- [51] **Business Standard:** [India must be part of AI safety regulation talks: IndiaAI COO](#) · 7 Oct 2026
- [52] **Washington Examiner:** [Top Senate Commerce Democrat's six-point AI framework](#) · 7 Oct 2026
- [53] **GitHub Changelog:** [Purpose-built model for leaked-secret detection](#) · 7 Oct 2026
- [54] **Python.org:** [Python 3.15.0rc3 release page](#) · 2 Oct 2026
- [55] **Java.com:** [Java release schedule](#) · Oct 2026
- [56] **Unite.AI:** [Automation Anywhere agrees to acquire Boost.ai](#) · 7 Oct 2026
- [57] **Benzinga:** [SpaceX, Apollo \\$40 billion Nvidia chip financing \(context...](#) · 6 Oct 2026
- [58] **CNBC:** [TSMC's September sales surge year over year](#) · 8 Oct 2026
- [59] **DIGITIMES:** [IC designers to pass cost increases through 2027](#) · 7 Oct 2026
- [60] **The Washington Times:** [Democrats compete to shape party's AI regulation stance](#) · 7 Oct 2026
- [61] **Zee Business:** [IMC 2026: next phase to focus on 6G, AI, chips](#) · 8 Oct 2026
- [62] **Kotak Neo:** [TCS Q2 results preview](#) · Oct 2026
- [63] **DealStreetAsia:** [Desible.ai, byteXL and others raise funding](#) · 7 Oct 2026
- [64] **Forbes:** [Judge strikes down latest \\$100,000 H-1B fee \(background: 1 Oct\)](#) · 1 Oct 2026
- [65] **Angel One:** [Tech Mahindra board to consider Q2 FY27 results on 15 Oct](#) · Oct 2026